


Markov Chain

(Andrey Markov)

- A random sequence in which the prob. for the next state only depends on the current state, not the history of states before that. (Markov property, memoryless)


$$[p(x_{i+1} | x_i, x_{i-1}, \dots) = p(x_{i+1} | x_i)]$$

Examples: "Drunkards walk" : move +1 or -1 with equal prob.




Probs. for next number dep. only on current number, not how the chain got there.

"Socks in duvet"




- Distribute walkers across states and iterate → evolve towards stationary / equil. populations across states
- Important concept: Ergodicity



A chain is ergodic if any state can be reached from any other state with a finite number of steps

Non-ergodic chain:



Markov Chain Monte Carlo

A method that produces a Markov chain of samples/events/steps whose stationary distribution is a given prob distribution

In other words: Generate a set of samples $x_i \sim p(x)$
by recording the steps of a Markov chain in x -space.

General procedure:

- Current state x_i
- Generate a new candidate x' using a proposal pdf that only depends on current state
- Apply acceptance rule. If accepted: $x_{i+1} = x'$
If rejected: $x_{i+1} = x_i$
- Repeat

For the list of samples $\{x_1, x_2, x_3, \dots\}$ to end up as a set of samples distributed according to $p(x)$, the acceptance rule should ensure

- Ergodicity (can visit all states)
- Detailed balance: $p(x_i) \text{Prob}(x_i \rightarrow x') = p(x') \text{Prob}(x' \rightarrow x_i)$
can have: \uparrow small \uparrow large $= \uparrow$ large \uparrow small

• The Metropolis (or Metropolis-Hastings) algorithm is one example that ensures this. $\left[\text{Basis: } \text{Prob}(x \rightarrow x') = T(x \rightarrow x') A(x \rightarrow x') \right]$

• Generate candidate according to some proposal prob. $T(x_i \rightarrow x')$

• Calculate acceptance prob.

$$A(x_i \rightarrow x') = \min\left(1, \frac{p(x')}{p(x_i)} \frac{T(x' \rightarrow x_i)}{T(x_i \rightarrow x')}\right)$$

If proposal is symmetric: $T(x' \rightarrow x_i) = T(x_i \rightarrow x')$ (Metropolis algo.)

$$A(x_i \rightarrow x') = \min\left(1, \frac{p(x')}{p(x_i)}\right)$$

• Generate random number r from $U(0,1)$

- If $r \leq A(x_i \rightarrow x')$, then accept: $x_{i+1} = x'$

- If $r > A(x_i \rightarrow x')$, then reject: $x_{i+1} = x_i$

• Note:

• Always accepts jump to state with higher prob. ($p(x') > p(x_i)$)

• Sometimes accepts jump to state with lower prob ($p(x') < p(x_i)$)
(with prob. that ensures detailed balance.)

$\left[\text{Wait long enough} \Rightarrow \text{all states should be explored (theoretically). Freedom in choosing proposal prob. dist.} \right]$

• Only depends on ratio $\frac{p(x')}{p(x)}$, so normalization constants (like the partition function) cancel! [Very important in proj. 4!]

• Suggested algorithm for Project 4 :

Repeat
N times
=
"one MC cycle"

↓
due to strong
correlations
between subseq.
samples in
the state space.

• Generate candidate state \bar{S}'

- Pick random spin from lattice
- Flip it!

• Find value of $\frac{p(\bar{S}')}{p(\bar{S}_i)}$ In an efficient way!
Don't evaluate $p(\bar{S}')$ and
 $p(\bar{S})$ separately, and
don't evaluate $\exp(\dots)$ all
the time!

• Generate random number $r \sim U(0,1)$

and carry out accept/reject step : accept if $r < \frac{p(\bar{S}')}{p(\bar{S}_i)}$

• Then use the new state \bar{S}_{i+1} to compute energy, magnetization, expectation values after this number of cycles, etc.

• Store relevant numbers

• Repeat

Two views on the MCMC in Project 4

① As model of (discretized) time evolution of system

- Each time step there are $\sim N$ random spin-environment interactions that can flip the spin. So the Markov chain

$$\bar{S}_0 \rightarrow \bar{S}_1 \rightarrow \bar{S}_2 \rightarrow \bar{S}_3 \rightarrow \dots$$

effectively means

$$\bar{S}(t_0) \rightarrow \bar{S}(t_1) \rightarrow \bar{S}(t_2) \rightarrow \bar{S}(t_3) \rightarrow \dots$$

- If system starts in state \bar{S}_0 that has low prob. under equilibrium assumption (Boltzmann), then letting it evolve will "equilibrate the system", i.e. take \bar{S} into regions of state space that are more likely in equilibrium.

↳ Connects to MCMC "burn-in"

↳ Some unfortunate terminology in the literature, which makes it sound like $p(\bar{S})$ is changing, when it's just the chain of \bar{S} samples that is converging towards a good approx to $p(\bar{S})$.

② As a method for sampling from $p(\bar{S})$

- Just some way to draw a set of samples $\bar{S}_1, \bar{S}_2, \bar{S}_3, \dots$ such that the distribution of samples (in the long run) corresponds to $p(\bar{S})$.
- Not directly related to the physics — equally applicable for any pdf, regardless of topic.
- The physics of equilibrium was already contained in our choice to use the Boltzmann dist. for $p(\bar{S})$

Burn-in

- Samples in MCMC are correlated, i.e. not "independent and identically dist. draws" (iid)

$$p(x_{i+1} | x_i) \neq p(x_{i+1})$$

- In the limit $n \rightarrow \infty$, we will have $\frac{n_{x \in [x, x+dx]}}{n} = p(x)dx$
↑
number of samples
- But we cannot do $n \rightarrow \infty$ in practice!
 \Rightarrow Our resulting distribution of samples (and hence our estimated expectation values) depends on starting point, which can be unreasonably improbable compared to the number of samples we draw!
- Usual solution:
 - Throw away first part of a chain of samples (burn-in)
 - Not use every new sample (even after burn-in)
 - Run many chains from different starting points and combine results.

